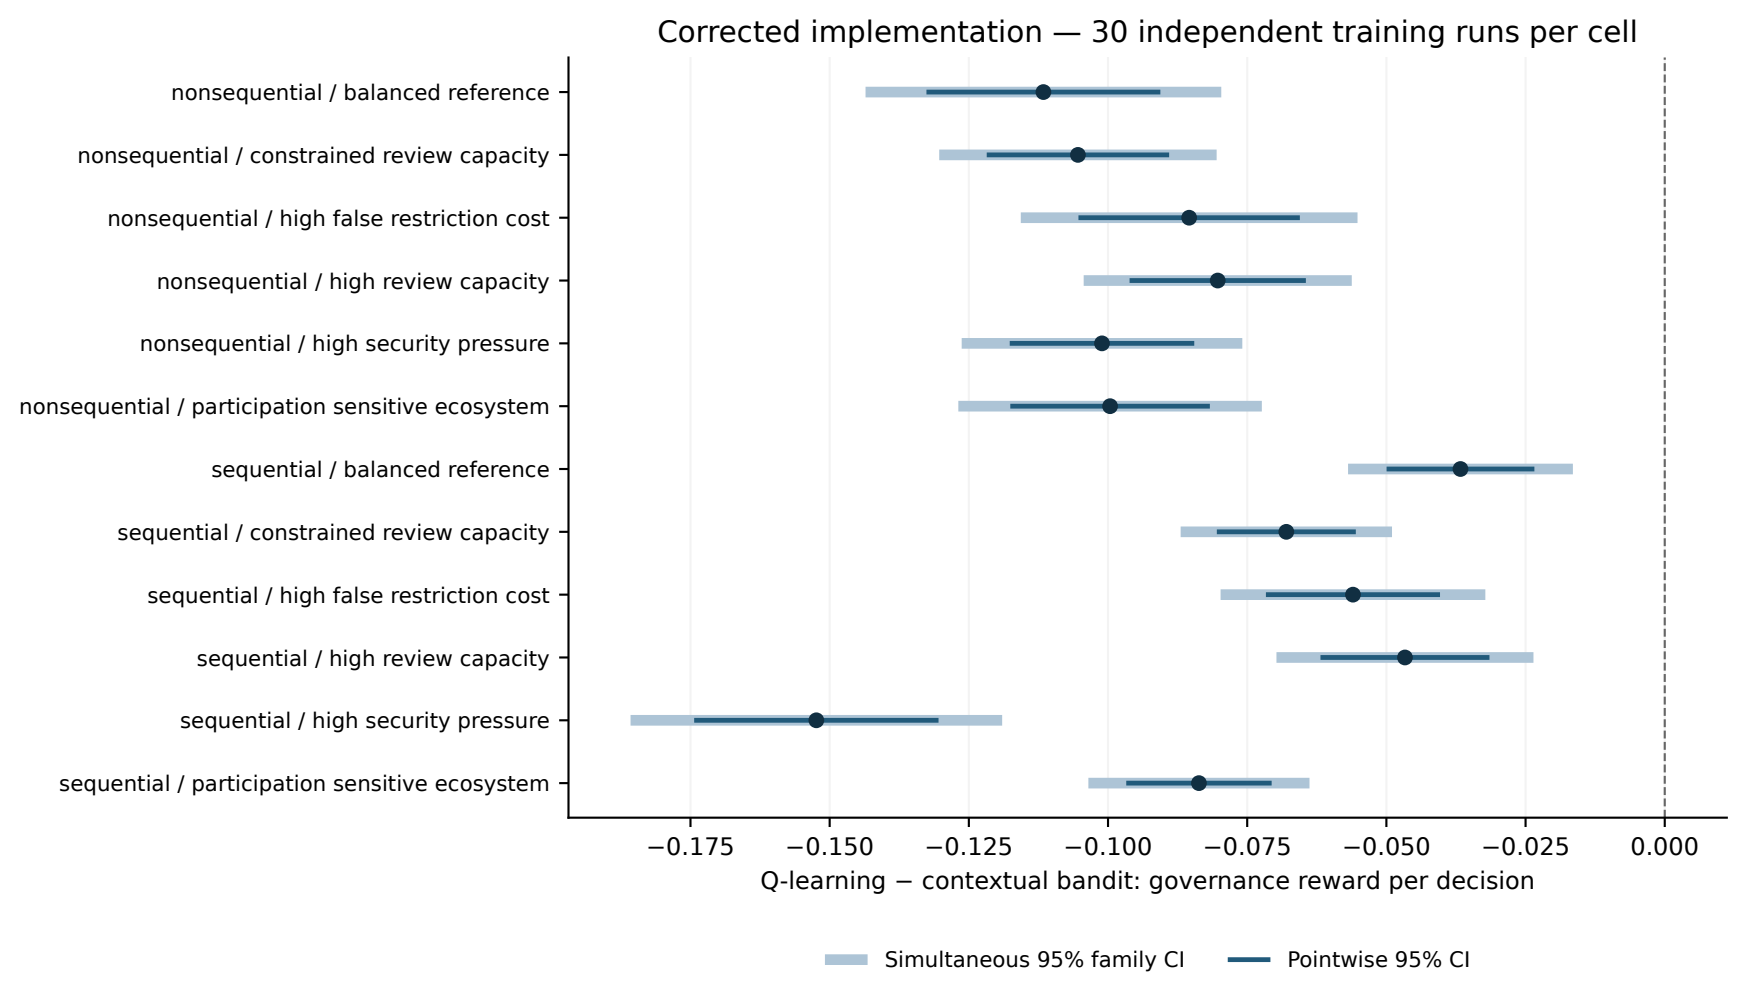
Corrected implementation — 30 independent training runs per cell
nonsequential / balanced reference
nonsequential / constrained review capacity
nonsequential / high false restriction cost
nonsequential / high review capacity
nonsequential / high security pressure
nonsequential / participation sensitive ecosystem
sequential / balanced reference
sequential / constrained review capacity
sequential / high false restriction cost
sequential / high review capacity
sequential / high security pressure
sequential / participation sensitive ecosystem
−0.175
−0.150
−0.125
−0.100
−0.075
−0.050
−0.025
0.000
Q-learning − contextual bandit: governance reward per decision
Simultaneous 95% family CI
Pointwise 95% CI